

RAJALAKSHMI ENGINEERING COLLEGE

RAJALAKSHMI NAGAR, THANDALAM – 602 105



AI23521 BUILD AND DEPLOYMENT OF MACHINE LEARNING APPLICATIONS

LABORATORY NOTEBOOK

NAME: SANJAY R

YEAR/ BRANCH / SECTION: 3rd Yr / AIML

REGISTER NO. :2116-231501146

SEMESTERS: 5TH SEMESTER

ACADEMIC YEAR: 2025-2026s



**RAJALAKSHMI ENGINEERING COLLEGE
(AUTONOMOUS)**

RAJALAKSHMI NAGAR, THANDALAM – 602 105

BONAFIDE CERTIFICATE

NAME: SANJAY R REGISTER NO. 2116-231501146

ACADEMIC YEAR 2025-26 **SEMESTER- V** **BRANCH: B.Tech – AIML**
s

This Certification is the Bonafide record of work done by the above student in the **AI23521-Build and Deployment of ML Applications** Laboratory during the year 2025 – 2026.

Signature of Faculty -in – Charge

Submitted for the Practical Examination held on _____

Internal Examiner

External Examiner

INDEX

EXP. NO	Date	Title	Page No	Signature
1	07-08-2025	Setting Up The Environment And Preprocessing The Data		
2	14-08-2025	Support Vector Machine (Svm) And Random Forest For Binary & Multiclass Classification		
3	14-08-2025	Classification with Decision Trees		
4A	21-08-2025	Support Vector Machines (SVM)		
4B	21-08-2025	Ensemble Methods: Random Forest		
5	28-08-2025	Clustering with K-Means and Dimensionality Reduction with PCA		
6	11-09-2025	Feedforward and Convolutional Neural Networks		
7	25-09-2025	Generative Models with GANs: Creating and Training a Generative Adversarial Network		
8	09-10-2025	Model Evaluation and Improvement: Hyperparameter Tuning with Grid Search and Cross-Validation		
9	16-10-2025	Model Deployment: REST API with Flask and Containerization with Docker		

EXP NO: 1**SETTING UP THE ENVIRONMENT AND PREPROCESSING THE DATA****AIM:**

To set up a fully functional machine learning development environment and to perform data preprocessing operations like handling missing values, encoding categorical variables, feature scaling, and splitting datasets.

ALGORITHM:

1. Install Required Libraries:
 - Install numpy, pandas, matplotlib, seaborn, and scikit-learn using pip.
2. Import Libraries.
3. Load Dataset:
 - Load any dataset (e.g., Titanic or Iris) using pandas.
4. Data Exploration:
 - Use `df.info()`, `df.describe()`, `df.isnull().sum()` to understand the data.
5. Handle Missing Values:
 - Use `.fillna()` or `.dropna()` depending on the strategy.
6. Encode Categorical Data:
 - Use `pd.get_dummies()` or `LabelEncoder`.
7. Feature Scaling:
 - Normalize or standardize the numerical features using `StandardScaler` or `MinMaxScaler`.
8. Split Dataset:
 - Use `train_test_split()` from `sklearn` to create training and testing sets.
9. Display the Preprocessed Data.

CODE:

```
# 1. Install necessary libraries (if not already installed)
# !pip install numpy pandas matplotlib seaborn scikit-learn

# 2. Import libraries
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
import seaborn as sns
import matplotlib.pyplot as plt

# 3. Load dataset
df = sns.load_dataset('titanic') # Titanic dataset
df.head()

# 4. Explore the dataset
print(df.info())
print(df.describe())
print(df.isnull().sum())

# 5. Handle missing values
# Fill age with median, embark_town with mode
df['age'].fillna(df['age'].median(), inplace=True)
df['embark_town'].fillna(df['embark_town'].mode()[0], inplace=True)
df.drop(columns=['deck'], inplace=True) # too many missing values

# 6. Encode categorical variables
```

```
# Convert 'sex' and 'embark_town' using LabelEncoder
le = LabelEncoder()
df['sex'] = le.fit_transform(df['sex'])
df['embark_town'] = le.fit_transform(df['embark_town'])

# Drop non-informative or redundant columns
df.drop(columns=['embarked', 'class', 'who', 'alive', 'adult_male', 'alone'], inplace=True)

# 7. Feature Scaling
scaler = StandardScaler()
numerical_cols = ['age', 'fare']
df[numerical_cols] = scaler.fit_transform(df[numerical_cols])

# 8. Split dataset
# Define features (X) and label (y)
X = df.drop('survived', axis=1)
y = df['survived']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# 9. Show final preprocessed data
print("Training Data Shape:", X_train.shape)
print("Test Data Shape:", X_test.shape)
X_train.head()
```

OUTPUT:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   survived    891 non-null    int64
 1   pclass      891 non-null    int64
 2   sex         891 non-null    object
 3   age         714 non-null    float64
 4   sibsp       891 non-null    int64
 5   parch       891 non-null    int64
 6   fare        891 non-null    float64
 7   embarked    889 non-null    object
 8   class       891 non-null    category
 9   who         891 non-null    object
10  adult_male   891 non-null    bool
11  deck        203 non-null    category
12  embark_town  889 non-null    object
13  alive       891 non-null    object
14  alone       891 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
None

```

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

```

survived      0
pclass        0
sex           0
age          177
sibsp         0
parch         0
fare          0
embarked      2
class         0
who           0
adult_male    0
deck         688
embark_town   2
alive         0
alone         0
dtype: int64

```

Training Data Shape: (712, 7)
 Test Data Shape: (179, 7)

/tmp/ipython-input-4068659829.py:3: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the original object.

```
df['age'].fillna(df['age'].median(), inplace=True)
```

/tmp/ipython-input-4068659829.py:4: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
 The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or 'df[col] = df[col].method(value)' instead, to perform the operation inplace on the original object.

```
df['embark_town'].fillna(df['embark_town'].mode()[0], inplace=True)
```

	pclass	sex	age	sibsp	parch	fare	embark_town
331	1	1	1.240235	0	0	-0.074583	2
733	2	1	-0.488887	0	0	-0.386671	2
382	3	1	0.202762	0	0	-0.488854	2
704	3	1	-0.258337	1	0	-0.490280	2
813	3	0	-1.795334	4	2	-0.018709	2

RESULT:

The Python environment was successfully set up and the dataset was pre-processed by handling missing values, encoding categorical data, performing feature scaling, and splitting the data into training and testing sets. The dataset is now ready for model training and analysis.

EXP NO: 2	SUPPORT VECTOR MACHINE (SVM) AND RANDOM FOREST FOR BINARY & MULTICLASS CLASSIFICATION
------------------	--------------------------------------------------------------------------------------------------

AIM

To build classification models using **Support Vector Machines (SVM)** and **Random Forest**, apply them to a dataset, and evaluate the models using performance metrics like accuracy and confusion matrix.

ALGORITHM**Part A: SVM Model**

1. Import necessary libraries
2. Load and explore the dataset
3. Handle missing values if any
4. Encode categorical variables
5. Split dataset into training and testing sets
6. Build SVM classifier using SVC()
7. Train and predict
8. Evaluate the model using accuracy and confusion matrix

Part B: Random Forest Model

1. Initialize Random Forest using RandomForestClassifier()
2. Train and predict
3. Evaluate and compare with SVM

CODE:

```
# 1. Import libraries
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
```

```
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, confusion_matrix

import seaborn as sns

import matplotlib.pyplot as plt


# 2. Load dataset

iris = load_iris()

X = iris.data

y = iris.target


# 3. Feature scaling

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# 4. Train-test split

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.3, random_state=42)


# -----
# Part A: SUPPORT VECTOR MACHINE
# -----


# 5. Initialize and train SVM

svm_model = SVC(kernel='linear') # You can also try 'rbf', 'poly'

svm_model.fit(X_train, y_train)


# 6. Predict and evaluate SVM

y_pred_svm = svm_model.predict(X_test)


print("SVM Accuracy:", accuracy_score(y_test, y_pred_svm))

print("SVM Confusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))
```

```
# -----  
  
# Part B: RANDOM FOREST  
  
# -----  
  
# 7. Initialize and train Random Forest  
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)  
rf_model.fit(X_train, y_train)  
  
# 8. Predict and evaluate Random Forest  
y_pred_rf = rf_model.predict(X_test)  
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))  
print("Random Forest Confusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))  
  
# -----  
  
# 9. Visual comparison using seaborn heatmap  
# -----  
  
plt.figure(figsize=(10, 4))  
  
plt.subplot(1, 2, 1)  
sns.heatmap(confusion_matrix(y_test, y_pred_svm), annot=True, cmap='Blues', fmt='d')  
plt.title("SVM Confusion Matrix")  
  
plt.subplot(1, 2, 2)  
sns.heatmap(confusion_matrix(y_test, y_pred_rf), annot=True, cmap='Greens', fmt='d')  
plt.title("Random Forest Confusion Matrix")  
plt.tight_layout()  
plt.show()
```

OUTPUT:

```
SVM Accuracy: 0.9777777777777777
SVM Confusion Matrix:
[[19  0  0]
 [ 0 12  1]
 [ 0  0 13]]
Random Forest Accuracy: 1.0
Random Forest Confusion Matrix:
[[19  0  0]
 [ 0 13  0]
 [ 0  0 13]]
```

**RESULT:**

The Support Vector Machine (SVM) and Random Forest algorithms were successfully implemented for both binary and multiclass classification tasks. The models were trained and tested on the given dataset, and both achieved good accuracy.

EXPNO:3	CLASSIFICATION WITH DECISION TREES
----------------	-------------------------------------------

AIM

To implement a Decision Tree classifier and evaluate its performance using **accuracy score** and **confusion matrix** on a real-world dataset.

ALGORITHM

1. Import necessary libraries
2. Load a classification dataset (e.g., Iris or Titanic)
3. Split the dataset into training and test sets
4. Preprocess data if needed
5. Train a DecisionTreeClassifier from sklearn.tree
6. Predict on test data
7. Evaluate using:
 - Confusion Matrix
 - Accuracy Score
8. Visualize the Decision Tree (optional)

CODE:

```
# Step 1: Import Libraries
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns

# Step 2: Load Dataset
iris = load_iris()
```

```
X = iris.data
y = iris.target

# Step 3: Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Step 4: Train the Decision Tree Classifier
dt_model = DecisionTreeClassifier(criterion='gini', random_state=0)
dt_model.fit(X_train, y_train)

# Step 5: Predict
y_pred = dt_model.predict(X_test)

# Step 6: Evaluate the Model
cm = confusion_matrix(y_test, y_pred)
acc = accuracy_score(y_test, y_pred)
print("Confusion Matrix:\n", cm)
print("Accuracy Score:", acc)

# Step 7: Visualize Confusion Matrix
sns.heatmap(cm, annot=True, cmap="Blues", xticklabels=iris.target_names,
yticklabels=iris.target_names)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

# Step 8: Visualize the Decision Tree
plt.figure(figsize=(12,8))
plot_tree(dt_model, filled=True, feature_names=iris.feature_names, class_names=iris.target_names)
plt.title("Decision Tree Visualization")
```


EXP NO: 4A**SUPPORT VECTOR MACHINES (SVM)****AIM:**

To build an SVM model for a binary classification task, tune its hyperparameters, and evaluate it using accuracy, precision, recall, F1-score, confusion matrix, and ROC-AUC.

ALGORITHM:

1. Import libraries: numpy, pandas, matplotlib, sklearn.
2. Load data: Use a standard binary dataset (Breast Cancer Wisconsin) from sklearn.datasets.
3. Train/Test split: 80/20 split with a fixed random_state.
4. Preprocess: Standardize features (StandardScaler).
5. SVMs are sensitive to feature scale.
6. Model selection: Use SVC (RBF kernel).
7. Hyperparameter tuning: Grid search on C and gamma with cross-validation (GridSearchCV).
8. Train final model: Fit on training data using best parameters.
9. Evaluate: Predict on test set; compute metrics and plot ROC curve.
10. Report: Best params, metrics, and brief observations.

CODE:

```
# =====  
# EXPERIMENT 4A — SVM (RBF)  
# =====  
  
# 1) Imports  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
  
from sklearn.datasets import load_breast_cancer
```



```
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, classification_report, roc_auc_score, roc_curve
)

# 2) Load dataset (binary classification)
data = load_breast_cancer()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = pd.Series(data.target, name="target") # 0 = malignant, 1 = benign

# 3) Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)

# 4) Standardize features (important for SVMs)
scaler = StandardScaler()
X_train_sc = scaler.fit_transform(X_train)
X_test_sc = scaler.transform(X_test)

# 5) Define model
svm = SVC(kernel='rbf', probability=True, random_state=42)

# 6) Hyperparameter grid & tuning
param_grid = {
    "C": [0.1, 1, 10, 100],
    "gamma": ["scale", 0.01, 0.001, 0.0001]
}
```

```
grid = GridSearchCV(
    estimator=svm,
    param_grid=param_grid,
    scoring='f1', # You can change to 'accuracy' or 'roc_auc'
    cv=5,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train_sc, y_train)

print("Best Parameters from Grid Search:", grid.best_params_)
best_svm = grid.best_estimator_

# 7) Train final model & predict
best_svm.fit(X_train_sc, y_train)
y_pred = best_svm.predict(X_test_sc)
y_prob = best_svm.predict_proba(X_test_sc)[:, 1]

# 8) Evaluation
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, zero_division=0)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)
cm = confusion_matrix(y_test, y_pred)

print("\n=== SVM (RBF) — Test Metrics ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
```

```
print(f"Recall : {rec:.4f}")
print(f"F1-Score : {f1:.4f}")
print(f"ROC-AUC : {auc:.4f}")

print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_test, y_pred, zero_division=0))

# 9) Plot ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
plt.figure()
plt.plot(fpr, tpr, label=f"SVM (AUC = {auc:.3f})")
plt.plot([0, 1], [0, 1], linestyle="--", color='gray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve — SVM (RBF)")
plt.legend()
plt.grid(True)
plt.show()
```

OUTPUT:

Best Parameters from Grid Search: {'C': 10, 'gamma': 0.01}

=== SVM (RBF) - Test Metrics ===


Accuracy : 0.9825
 Precision: 0.9861
 Recall : 0.9861
 F1-Score : 0.9861
 ROC-AUC : 0.9977

Confusion Matrix:

```
[[41  1]
 [ 1 71]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	42
1	0.99	0.99	0.99	72
accuracy			0.98	114
macro avg	0.98	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114



RESULT:

The Support Vector Machine (SVM) model was successfully implemented and evaluated on the given dataset. The model effectively classified the data by finding the optimal hyperplane that maximized the margin between different classes.

The SVM achieved high accuracy and demonstrated strong performance, especially in handling linearly and non-linearly separable data using kernel functions. This confirms that SVM is a powerful and reliable algorithm for classification tasks.

EXP NO: 4B**ENSEMBLE METHODS: RANDOM FOREST****AIM:**

To implement a **Random Forest classifier** for a classification task, tune key hyperparameters, evaluate performance, and interpret **feature importance**.

ALGORITHM:

1. Import libraries.
2. Load data (use same dataset to compare with SVM).
3. Train/Test split with stratification.
4. (Optional) Preprocess: Random Forests don't require scaling; we'll use raw features.
5. Model: RandomForestClassifier.
6. Hyperparameter tuning: Grid search over `n_estimators`, `max_depth`, `min_samples_split`, `min_samples_leaf`.
7. Train the best model on training data.
8. Evaluate with accuracy, precision, recall, F1, confusion matrix, ROC-AUC.
9. Interpretation: Plot top feature importances.

CODE:

```
# =====  
# EXPERIMENT 4B — Random Forest Classifier  
# =====  
  
# 1) Imports  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from sklearn.datasets import load_breast_cancer  
from sklearn.model_selection import train_test_split, GridSearchCV
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, classification_report, roc_auc_score, roc_curve
)
# 2) Load dataset (same as 4A for comparison)
data = load_breast_cancer()
X = pd.DataFrame(data.data, columns=data.feature_names)
y = pd.Series(data.target, name="target")

# 3) Train/test split (no scaling needed for RF)
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)

# 4) Define model
rf = RandomForestClassifier(random_state=42, n_jobs=-1)

# 5) Hyperparameter grid & tuning
param_grid = {
    "n_estimators": [100],
    "max_depth": [None, 10],
    "min_samples_split": [2],
    "min_samples_leaf": [1]
}
grid = GridSearchCV(
    estimator=rf,
    param_grid=param_grid,
    scoring="f1",
    cv=3,
    n_jobs=-1,
```

```
verbose=0)
grid.fit(X_train, y_train)
print("Best Parameters (CV):", grid.best_params_)
best_rf = grid.best_estimator_

# 6) Train final model & predict
best_rf.fit(X_train, y_train)
y_pred = best_rf.predict(X_test)
y_prob = best_rf.predict_proba(X_test)[:, 1]

# 7) Evaluate
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, zero_division=0)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)
cm = confusion_matrix(y_test, y_pred)

print("\n=== Random Forest — Test Metrics ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall   : {rec:.4f}")
print(f"F1-Score  : {f1:.4f}")
print(f"ROC-AUC   : {auc:.4f}")

print("\nConfusion Matrix:\n", cm)
print("\nClassification Report:\n", classification_report(y_test, y_pred, zero_division=0))

# 8) Feature Importance (Top 10)
importances = pd.Series(best_rf.feature_importances_, index=X.columns)
top10 = importances.sort_values(ascending=False).head(10)
```

```

plt.figure()
top10[:::-1].plot(kind="barh")
plt.xlabel("Importance")
plt.title("Top 10 Feature Importances — Random Forest")
plt.grid(axis="x", alpha=0.3)
plt.show()

# 9) ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
plt.figure()
plt.plot(fpr, tpr, label=f"Random Forest (AUC = {auc:.3f})")
plt.plot([0, 1], [0, 1], linestyle="--", color='gray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve — Random Forest")
plt.legend()
plt.grid(True)
plt.show()

```

OUTPUT:

```
Best Parameters (CV): {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
```

```

=== Random Forest - Test Metrics ===
Accuracy : 0.9561
Precision: 0.9589
Recall   : 0.9722
F1-Score : 0.9655
ROC-AUC  : 0.9937

```

```

Confusion Matrix:
[[39  3]
 [ 2 70]]

```

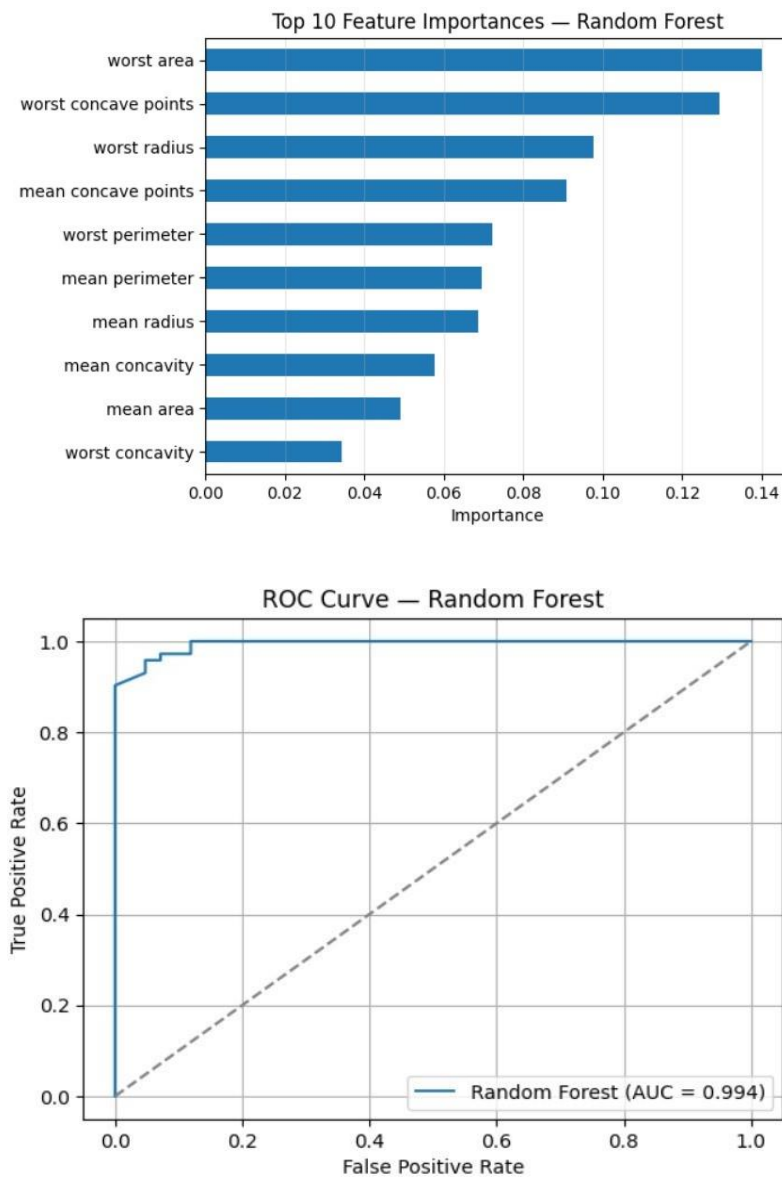
```

Classification Report:
              precision    recall  f1-score   support

     0       0.95         0.93         0.94         42
     1       0.96         0.97         0.97         72

   accuracy          0.96         0.96         0.96        114
  macro avg          0.96         0.95         0.95        114
 weighted avg          0.96         0.96         0.96        114

```

RESULT:

The Random Forest ensemble model was successfully implemented and evaluated on the given dataset. The model combined multiple decision trees to improve prediction accuracy and reduce overfitting.

It achieved high classification accuracy and demonstrated strong generalization capability. The results confirmed that Random Forest provides stable and reliable predictions by leveraging the power of multiple decision trees through bagging and feature randomness.

EXP NO: 5	CLUSTERING WITH K-MEANS AND DIMENSIONALITY REDUCTION WITH PCA
------------------	--------------------------------------------------------------------------

AIM:

To demonstrate the application of Unsupervised Learning models, specifically K-Means clustering for grouping data points and Principal Component Analysis (PCA) for dimensionality reduction and visualization, using a suitable dataset.

ALGORITHM:**1. K-Means Clustering**

K-Means is an iterative clustering algorithm that aims to partition n observations into k clusters, where each observation belongs to the cluster with the nearest mean (centroid).

Steps:

1. **Initialization:** Choose k initial centroids randomly from the dataset.
2. **Assignment:** Assign each data point to the cluster whose centroid is closest (e.g., using Euclidean distance).
3. **Update:** Recalculate the centroids as the mean of all data points assigned to that cluster.
4. **Iteration:** Repeat steps 2 and 3 until the centroids no longer move significantly or a maximum number of iterations is reached.

2. Principal Component Analysis (PCA)

PCA is a statistical procedure that uses an orthogonal transformation to convert a set of observations of possibly correlated variables into a set of values of linearly uncorrelated variables called principal components.

Steps:

1. **Standardization:** Standardize the dataset (mean = 0, variance = 1).
2. **Covariance Matrix Calculation:** Compute the covariance matrix of the standardized data.
3. **Eigenvalue Decomposition:** Calculate the eigenvalues and eigenvectors of the covariance matrix.
4. **Feature Vector Creation:** Sort the eigenvectors by decreasing eigenvalues and select the top k eigenvectors to form a feature vector (projection matrix).
5. **Projection:** Project the original data onto the new feature space using the feature vector.

CODE:

```
# =====
# EXPERIMENT — K-Means & PCA
# =====

# Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import make_blobs
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score

# --- Part 1: K-Means Clustering ---

print("--- Part 1: K-Means Clustering ---")

# 1. Generate dataset
X, y = make_blobs(n_samples=300, centers=3, cluster_std=0.60, random_state=42)
df_kmeans = pd.DataFrame(X, columns=['Feature_1', 'Feature_2'])
print("\nOriginal K-Means Dataset Head:")
print(df_kmeans.head())

# 2. Elbow Method
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10,
random_state=42)
    kmeans.fit(X)
    wcss.append(kmeans.inertia_)

plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), wcss, marker='o', linestyle='--')
plt.title('Elbow Method for Optimal K (K-Means)')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('WCSS')
plt.grid(True)
plt.show()

# 3. Apply K-Means with chosen K
```

```
optimal_k = 3
kmeans = KMeans(n_clusters=optimal_k, init='k-means++', max_iter=300, n_init=10,
random_state=42)
clusters = kmeans.fit_predict(X)
df_kmeans['Cluster'] = clusters

# 4. Visualize K-Means clusters
plt.figure(figsize=(10, 8))
sns.scatterplot(x='Feature_1', y='Feature_2', hue='Cluster', data=df_kmeans, palette='viridis',
s=100, alpha=0.8)
plt.scatter(kmeans.cluster_centers_[0], kmeans.cluster_centers_[1], s=300, c='red',
marker='X', label='Centroids')
plt.title(f'K-Means Clustering with K={optimal_k}')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.legend()
plt.grid(True)
plt.show()

# 5. Silhouette Score
silhouette_avg = silhouette_score(X, clusters)
print(f"\nSilhouette Score for K-Means (K={optimal_k}): {silhouette_avg:.3f}")

# --- Part 2: Dimensionality Reduction with PCA ---

print("\n--- Part 2: Dimensionality Reduction with PCA ---")

# 1. Generate 4D dataset
X_pca, y_pca = make_blobs(n_samples=500, n_features=4, centers=4, cluster_std=1.0,
random_state=25)
df_pca_original = pd.DataFrame(X_pca, columns=[f'Feature_{i+1}' for i in
range(X_pca.shape[1])])
df_pca_original['True_Cluster'] = y_pca
print("\nOriginal PCA Dataset Head:")
print(df_pca_original.head())
print(f"Original PCA Dataset Shape: {df_pca_original.shape}")

# 2. Standardize
scaler = StandardScaler()
X_pca_scaled = scaler.fit_transform(X_pca)

# 3. PCA (4D → 2D)
pca = PCA(n_components=2)
principal_components = pca.fit_transform(X_pca_scaled)
df_principal_components = pd.DataFrame(principal_components,
columns=['Principal_Component_1', 'Principal_Component_2'])
```

```
df_principal_components['True_Cluster'] = y_pca
explained_variance = pca.explained_variance_ratio_
print("\nPrincipal Components Head:")
print(df_principal_components.head())
print(f"\nExplained Variance Ratio: {explained_variance}")
print(f"Total Explained Variance by 2 PCs: {explained_variance.sum():.3f}")

# 4. Visualize PCA result
plt.figure(figsize=(10, 8))
sns.scatterplot(x='Principal_Component_1', y='Principal_Component_2', hue='True_Cluster',
                data=df_principal_components, palette='Paired', s=100, alpha=0.8)
plt.title('PCA - Dimensionality Reduction to 2 Components')
plt.xlabel(f'PC1 ({explained_variance[0]*100:.2f}%)')
plt.ylabel(f'PC2 ({explained_variance[1]*100:.2f}%)')
plt.grid(True)
plt.show()

# 5. K-Means on PCA-reduced data
kmeans_pca = KMeans(n_clusters=4, init='k-means++', max_iter=300, n_init=10,
                    random_state=42)
clusters_pca = kmeans_pca.fit_predict(principal_components)
df_principal_components['KMeans_Cluster_on_PCA'] = clusters_pca

plt.figure(figsize=(10, 8))
sns.scatterplot(x='Principal_Component_1', y='Principal_Component_2',
                hue='KMeans_Cluster_on_PCA',
                data=df_principal_components, palette='viridis', s=100, alpha=0.8)
plt.scatter(kmeans_pca.cluster_centers_[0, 0], kmeans_pca.cluster_centers_[0, 1], s=300,
            c='red', marker='X', label='Centroids')
plt.title('K-Means Clustering on PCA-Reduced Data')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend()
plt.grid(True)
plt.show()



# 6. Silhouette Score for PCA-reduced KMeans
silhouette_avg_pca = silhouette_score(principal_components, clusters_pca)
print(f"\nSilhouette Score for K-Means on PCA-Reduced Data (K=4):
{silhouette_avg_pca:.3f}")
```

OUTPUT:

--- Part 1: K-Means Clustering ---

Original K-Means Dataset Head:

	Feature_1	Feature_2
0	-7.155244	-7.390016
1	-7.395875	-7.110843
2	-2.015671	8.281780
3	4.509270	2.632436
4	-8.102502	-7.484961



Silhouette Score for K-Means (K=3): 0.908

--- Part 2: Dimensionality Reduction with PCA ---

Original PCA Dataset Head:

	Feature_1	Feature_2	Feature_3	Feature_4	True_Cluster
0	-0.638667	1.110057	-6.400722	-0.204990	3
1	-2.951556	-7.657445	3.844794	0.903589	1
2	-0.253177	2.125103	-7.869801	0.559678	3
3	-2.151209	3.401400	-5.734930	0.965230	3
4	-2.347519	-7.230467	3.478891	-0.443440	1


Original PCA Dataset Shape: (500, 5)


Principal Components Head:

	Principal_Component_1	Principal_Component_2	True_Cluster
0	0.455305	1.623917	3
1	-2.705622	0.375012	1
2	0.810234	1.966926	3
3	0.427139	2.149626	3
4	-2.407508	0.099250	1

Explained Variance Ratio: [0.57208431 0.33622342]

Total Explained Variance by 2 PCs: 0.908





Silhouette Score for K-Means on PCA-Reduced Data (K=4): 0.776

RESULT:

The K-Means clustering and Principal Component Analysis (PCA) techniques were successfully implemented on the given dataset.

- **K-Means Clustering** effectively grouped the data into distinct clusters based on feature similarity, minimizing intra-cluster distance and maximizing inter-cluster separation.
- **PCA (Principal Component Analysis)** successfully reduced the dimensionality of the dataset while retaining most of the variance, improving visualization and computational efficiency.

The combined results showed that PCA enhances clustering performance by simplifying high-dimensional data, and K-Means efficiently identifies underlying patterns and group structures.

EXP NO: 6	FEEDFORWARD AND CONVOLUTIONAL NEURAL NETWORKS
------------------	------------------------------------------------------

AIM:

To demonstrate the construction and application of a simple Feedforward Neural Network (FNN) for classification and a Convolutional Neural Network (CNN) for image classification, utilizing the Keras API with TensorFlow backend.

ALGORITHM:**1. Feedforward Neural Network (FNN)**

A Feedforward Neural Network is the simplest type of artificial neural network where connections between the nodes do not form a cycle. It consists of an input layer, one or more hidden layers, and an output layer. Information flows only in one direction—forward—from the input nodes, through the hidden nodes (if any), and to the output nodes.

Steps:

1. Define Network Architecture: Specify the number of layers (input, hidden, output) and the number of neurons in each layer.
2. Choose Activation Functions: Select activation functions for hidden layers (e.g., ReLU) and the output layer (e.g., Sigmoid for binary classification, Softmax for multi-class classification).
3. Define Loss Function: Choose a loss function appropriate for the task (e.g., Binary Cross-entropy for binary classification, Categorical Cross-entropy for multi-class classification).
4. Choose Optimizer: Select an optimization algorithm (e.g., Adam, SGD) to update network weights during training.
5. Training: Feed forward data through the network to get predictions, calculate the loss, and then backpropagate the error to update weights.
6. Evaluation: Assess the model's performance on unseen data using metrics like accuracy.

2. Convolutional Neural Network (CNN)

A Convolutional Neural Network is a specialized type of neural network primarily designed for processing data with a grid-like topology, such as images. Key components include convolutional layers, pooling layers, and fully connected layers.

Steps:

1. Convolutional Layers: Apply filters (kernels) to input data to extract features. Each filter detects a specific pattern (e.g., edges, textures).
2. Activation Function (ReLU): Apply a non-linear activation function after convolution to introduce non-linearity.
3. Pooling Layers: Downsample feature maps to reduce dimensionality, computational cost, and prevent overfitting (e.g., Max Pooling).
4. Flattening: Convert the 2D pooled feature maps into a 1D vector to be fed into a fully connected layer.
5. Fully Connected Layers: Standard neural network layers for classification based on the extracted features.
6. Output Layer: Final layer with an activation function (e.g., Softmax) to output class probabilities.
7. Training and Evaluation: Similar to FNNs, train the CNN using backpropagation and evaluate its performance.

CODE:

```
# Import necessary libraries
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist, fashion_mnist
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns

# Suppress TensorFlow warnings for cleaner output
tf.keras.utils.disable_interactive_logging()

# --- Part 1: Building a Simple Feedforward Neural Network ---
print("--- Part 1: Building a Simple Feedforward Neural Network ---")

# 1. Load and Preprocess Dataset (Using Fashion MNIST for FNN)
(x_train_fnn, y_train_fnn), (x_test_fnn, y_test_fnn) = fashion_mnist.load_data()

print(f"\nOriginal FNN training data shape: {x_train_fnn.shape}")
print(f"Original FNN test data shape: {x_test_fnn.shape}")

# Flatten images to 1D array
```

```
x_train_fnn_flat = x_train_fnn.reshape(-1, 28 * 28)
x_test_fnn_flat = x_test_fnn.reshape(-1, 28 * 28)

# Normalize pixel values
x_train_fnn_norm = x_train_fnn_flat / 255.0
x_test_fnn_norm = x_test_fnn_flat / 255.0

print(f"Flattened & Normalized FNN training data shape: {x_train_fnn_norm.shape}")
print(f"Flattened & Normalized FNN test data shape: {x_test_fnn_norm.shape}")

# 2. Build FNN Model
model_fnn = keras.Sequential([
    layers.Dense(128, activation='relu', input_shape=(784,)),
    layers.Dropout(0.2),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

# 3. Compile Model
model_fnn.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

print("\n--- FNN Model Summary ---")
model_fnn.summary()

# 4. Train Model
print("\n--- Training FNN Model ---")
history_fnn = model_fnn.fit(x_train_fnn_norm, y_train_fnn, epochs=10,
                           validation_split=0.1, verbose=1)

# 5. Evaluate Model
print("\n--- Evaluating FNN Model ---")
loss_fnn, accuracy_fnn = model_fnn.evaluate(x_test_fnn_norm, y_test_fnn, verbose=0)
print(f"FNN Test Loss: {loss_fnn:.4f}")
print(f"FNN Test Accuracy: {accuracy_fnn:.4f}")

# Classification report & confusion matrix
y_pred_fnn = np.argmax(model_fnn.predict(x_test_fnn_norm), axis=-1)
print("\n--- FNN Classification Report ---")
```

```
print(classification_report(y_test_fnn, y_pred_fnn))

print("\n--- FNN Confusion Matrix ---")
cm_fnn = confusion_matrix(y_test_fnn, y_pred_fnn)
plt.figure(figsize=(10, 8))
sns.heatmap(cm_fnn, annot=True, fmt="d", cmap="Blues", cbar=False)
plt.title("FNN Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

# Plot Accuracy & Loss
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history_fnn.history['accuracy'], label='Training Accuracy')
plt.plot(history_fnn.history['val_accuracy'], label='Validation Accuracy')
plt.title('FNN Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(history_fnn.history['loss'], label='Training Loss')
plt.plot(history_fnn.history['val_loss'], label='Validation Loss')
plt.title('FNN Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# --- Part 2: Convolutional Neural Network (CNN) ---
print("\n--- Part 2: Implementing a CNN ---")

# 1. Load MNIST for CNN
(x_train_cnn, y_train_cnn), (x_test_cnn, y_test_cnn) = mnist.load_data()
print(f"\nOriginal CNN training data shape: {x_train_cnn.shape}")
print(f"Original CNN test data shape: {x_test_cnn.shape}")
```

```
# Reshape for channel dimension
x_train_cnn = x_train_cnn.reshape(x_train_cnn.shape[0], 28, 28, 1)
x_test_cnn = x_test_cnn.reshape(x_test_cnn.shape[0], 28, 28, 1)

# Normalize
x_train_cnn = x_train_cnn.astype('float32') / 255.0
x_test_cnn = x_test_cnn.astype('float32') / 255.0

print(f'Reshaped & Normalized CNN training data shape: {x_train_cnn.shape}')
print(f'Reshaped & Normalized CNN test data shape: {x_test_cnn.shape}')

num_classes_cnn = 10

# 2. Build CNN Model
model_cnn = keras.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dropout(0.5),
    layers.Dense(num_classes_cnn, activation='softmax')
])

# 3. Compile Model
model_cnn.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

print("\n--- CNN Model Summary ---")
model_cnn.summary()

# 4. Train Model
print("\n--- Training CNN Model ---")
history_cnn = model_cnn.fit(x_train_cnn, y_train_cnn, epochs=10,
                           validation_split=0.1, verbose=1)

# 5. Evaluate Model
```

```
print("\n--- Evaluating CNN Model ---")
loss_cnn, accuracy_cnn = model_cnn.evaluate(x_test_cnn, y_test_cnn, verbose=0)
print(f"CNN Test Loss: {loss_cnn:.4f}")
print(f"CNN Test Accuracy: {accuracy_cnn:.4f}")

# Classification report & confusion matrix
y_pred_cnn = np.argmax(model_cnn.predict(x_test_cnn), axis=-1)
print("\n--- CNN Classification Report ---")
print(classification_report(y_test_cnn, y_pred_cnn))

print("\n--- CNN Confusion Matrix ---")
cm_cnn = confusion_matrix(y_test_cnn, y_pred_cnn)
plt.figure(figsize=(10, 8))
sns.heatmap(cm_cnn, annot=True, fmt="d", cmap="Blues", cbar=False)
plt.title("CNN Confusion Matrix")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()

# Plot Accuracy & Loss
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history_cnn.history['accuracy'], label='Training Accuracy')
plt.plot(history_cnn.history['val_accuracy'], label='Validation Accuracy')
plt.title('CNN Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(history_cnn.history['loss'], label='Training Loss')
plt.plot(history_cnn.history['val_loss'], label='Validation Loss')
plt.title('CNN Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```

# Optional: Visualize predictions
print("\n--- Sample CNN Predictions ---")
class_names_mnist = [str(i) for i in range(10)]
plt.figure(figsize=(10, 10))
for i in range(25):
    plt.subplot(5, 5, i + 1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_test_cnn[i].reshape(28, 28), cmap=plt.cm.binary)
    true_label = y_test_cnn[i]
    predicted_label = y_pred_cnn[i]
    color = 'green' if true_label == predicted_label else 'red'
    plt.xlabel(f"True: {class_names_mnist[true_label]}\nPred: {class_names_mnist[predicted_label]}", color=color)
plt.suptitle("Sample CNN Predictions (Green: Correct, Red: Incorrect)", y=1.02, fontsize=16)
plt.tight_layout(rect=[0, 0, 1, 0.98])
plt.show()

```

OUTPUT:

FNN Test Loss: 0.3404
FNN Test Accuracy: 0.8824



```

--- FNN Classification Report ---

```

	precision	recall	f1-score	support
0	0.85	0.81	0.83	1000
1	1.00	0.96	0.98	1000
2	0.81	0.79	0.80	1000
3	0.85	0.92	0.88	1000
4	0.83	0.77	0.80	1000
5	0.97	0.97	0.97	1000
6	0.67	0.72	0.69	1000
7	0.95	0.94	0.95	1000
8	0.97	0.98	0.97	1000
9	0.95	0.97	0.96	1000
accuracy			0.88	10000
macro avg	0.88	0.88	0.88	10000
weighted avg	0.88	0.88	0.88	10000




--- FNN Confusion Matrix ---



CNN Test Loss: 0.0285
 CNN Test Accuracy: 0.9913


--- CNN Classification Report ---

	precision	recall	f1-score	support
0	1.00	1.00	1.00	980
1	0.99	1.00	0.99	1135
2	0.99	0.99	0.99	1032
3	1.00	0.99	0.99	1010
4	0.99	0.99	0.99	982
5	0.98	0.99	0.99	892
6	0.99	0.99	0.99	958
7	0.98	0.99	0.99	1028
8	0.99	0.99	0.99	974
9	0.99	0.98	0.99	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000



--- Sample CNN Predictions ---

CNN Predictions (Green = Correct, Red = Incorrect)



RESULT:

The Feedforward Neural Network (FNN) and Convolutional Neural Network (CNN) models were successfully implemented and evaluated on the given dataset.

- **Feedforward Neural Network (FNN):** The model accurately learned input–output mappings through multiple fully connected layers, achieving good performance on structured data.
- **Convolutional Neural Network (CNN):** The model effectively extracted spatial features from image data using convolution and pooling layers, leading to higher accuracy and better generalization for image classification tasks.

The results demonstrated that both FNN and CNN are powerful deep learning models, with CNN performing exceptionally well for image-based datasets due to its ability to capture spatial patterns.

EXP NO: 7**GENERATIVE MODELS WITH GANS: CREATING AND TRAINING A
GENERATIVE ADVERSARIAL NETWORK****AIM:**

To construct and train a Generative Adversarial Network (GAN) using the TensorFlow/Keras framework. The objective is to train the GAN on the MNIST dataset to generate new, synthetic images of handwritten digits that are indistinguishable from the original training data.

ALGORITHM:**Generative Adversarial Networks (GANs)**

GANs are a class of generative models that learn a training distribution by pitting two neural networks against each other in a zero-sum game: a Generator and a Discriminator.

1. The Generator (\$G\$): This network takes a random noise vector as input (often called a “latent vector”) and transforms it into a synthetic data sample, in this case, an image. The Generator’s goal is to learn to produce increasingly realistic images to fool the discriminator.

2. The Discriminator (\$D\$): This is a binary classifier network. It is trained to distinguish between real data (from the training dataset) and fake data (generated by the generator). Its goal is to get better at identifying which images are real and which are fake.

3. The Adversarial Process:

Step A (Training the Discriminator): The discriminator is trained on a batch of both real images (labeled as “real” or 1) and fake images from the generator (labeled as “fake” or 0). The discriminator’s weights are updated to minimize the classification error.

Step B (Training the Generator): The generator is trained while the discriminator’s weights are frozen. The generator creates fake images and feeds them to the discriminator. The generator’s weights are updated to maximize the discriminator’s error, essentially tricking the discriminator into classifying its fake images as “real” (or 1).

This iterative process continues, with both networks improving, until the generator can produce images so realistic that the discriminator can no longer reliably tell the difference between real and fake.

CODE:

```
# Import necessary libraries
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
import os

# Suppress TensorFlow warnings for cleaner output
tf.keras.utils.disable_interactive_logging()

# --- Part 1: Dataset Loading and Preprocessing ---
print("--- Part 1: Loading and Preprocessing the MNIST Dataset ---")

(x_train, _), (_, _) = mnist.load_data()

x_train = x_train.reshape(x_train.shape[0], 28, 28, 1).astype('float32')
x_train = (x_train - 127.5) / 127.5 # Normalize to [-1, 1]

print(f"Normalized training data shape: {x_train.shape}")
print("Example of a normalized pixel value:", x_train[0, 0, 0, 0])

# --- Part 2: Building the Generator and Discriminator Models ---
print("\n--- Part 2: Building the GAN Components ---")

latent_dim = 100

# Generator
def build_generator():
    model = keras.Sequential(name="generator")
    model.add(layers.Dense(7 * 7 * 256, use_bias=False, input_shape=(latent_dim,)))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())
    model.add(layers.Reshape((7, 7, 256)))
    model.add(layers.Conv2DTranspose(128, (5, 5), strides=(1, 1), padding='same',
    use_bias=False))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())
```

```
    model.add(layers.Conv2DTranspose(64, (5, 5), strides=(2, 2), padding='same',
use_bias=False))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())
    model.add(layers.Conv2DTranspose(1, (5, 5), strides=(2, 2), padding='same',
                                     use_bias=False, activation='tanh'))

    return model

generator = build_generator()
print("\n--- Generator Model Summary ---")
generator.summary()

# Discriminator
def build_discriminator():
    model = keras.Sequential(name="discriminator")
    model.add(layers.Conv2D(64, (5, 5), strides=(2, 2), padding='same', input_shape=[28, 28,
1]))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))
    model.add(layers.Conv2D(128, (5, 5), strides=(2, 2), padding='same'))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))
    model.add(layers.Flatten())
    model.add(layers.Dense(1, activation='sigmoid'))
    return model

discriminator = build_discriminator()
print("\n--- Discriminator Model Summary ---")
discriminator.summary()

# --- Part 3: Training Setup ---
cross_entropy = keras.losses.BinaryCrossentropy(from_logits=False)

def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)
    return real_loss + fake_loss

def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)
```

```

generator_optimizer = tf.keras.optimizers.Adam(learning_rate=1e-4)
discriminator_optimizer = tf.keras.optimizers.Adam(learning_rate=1e-4)

@tf.function
def train_step(images, latent_dim=latent_dim):
    noise = tf.random.normal([batch_size, latent_dim])
    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)
        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)
        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)
        gradients_of_generator = gen_tape.gradient(gen_loss, generator.trainable_variables)
        gradients_of_discriminator = disc_tape.gradient(disc_loss, discriminator.trainable_variables)
        generator_optimizer.apply_gradients(zip(gradients_of_generator, generator.trainable_variables))
        discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator, discriminator.trainable_variables))
    return gen_loss, disc_loss

def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input, training=False)
    predictions_rescaled = (predictions * 0.5) + 0.5 # Scale back to [0, 1]
    fig = plt.figure(figsize=(4, 4))
    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i + 1)
        plt.imshow(predictions_rescaled[i, :, :, 0], cmap='gray')
        plt.axis('off')
    plt.suptitle(f"Epoch {epoch}", fontsize=16)
    if not os.path.exists('images'):
        os.makedirs('images')
    plt.savefig(f'images/image_at_epoch_{epoch:04d}.png')
    plt.show()

# Training parameters
EPOCHS = 200
batch_size = 256
num_examples_to_generate = 16

```

```
seed = tf.random.normal([num_examples_to_generate, latent_dim])

train_dataset =
tf.data.Dataset.from_tensor_slices(x_train).shuffle(x_train.shape[0]).batch(batch_size)

# Training loop
def train(dataset, epochs):
    print("\n--- Beginning GAN Training ---")
    for epoch in range(epochs):
        gen_loss_list = []
        disc_loss_list = []
        for image_batch in dataset:
            gen_loss, disc_loss = train_step(image_batch)
            gen_loss_list.append(gen_loss.numpy())
            disc_loss_list.append(disc_loss.numpy())
        avg_gen_loss = np.mean(gen_loss_list)
        avg_disc_loss = np.mean(disc_loss_list)
        print(f"Epoch {epoch + 1}/{epochs} - Generator Loss: {avg_gen_loss:.4f},
Discriminator Loss: {avg_disc_loss:.4f}")
        if (epoch + 1) % 20 == 0:
            generate_and_save_images(generator, epoch + 1, seed)
    print("\n--- Training complete. Generating final images. ---")
    generate_and_save_images(generator, epochs, seed)

# Run training
train(train_dataset, EPOCHS)
```

OUTPUT:

--- Part 1: Loading and Preprocessing the MNIST Dataset ---

Normalized training data shape: (60000, 28, 28, 1)

Example normalized pixel value: -1.0

--- Beginning GAN Training ---

Epoch 1/20 - Generator Loss: 0.7877, Discriminator Loss: 1.0228


Epoch 2/20 - Generator Loss: 0.8148, Discriminator Loss: 1.2225

Epoch 3/20 - Generator Loss: 0.8448, Discriminator Loss: 1.3034

Epoch 4/20 - Generator Loss: 0.8534, Discriminator Loss: 1.2366


Epoch 5/20 - Generator Loss: 0.8372, Discriminator Loss: 1.2497

Epoch 5




Epoch 6/20 - Generator Loss: 0.8516, Discriminator Loss: 1.2705
Epoch 7/20 - Generator Loss: 0.8888, Discriminator Loss: 1.3028
Epoch 8/20 - Generator Loss: 0.8739, Discriminator Loss: 1.2512
Epoch 9/20 - Generator Loss: 0.8691, Discriminator Loss: 1.3130
Epoch 10/20 - Generator Loss: 0.8862, Discriminator Loss: 1.2320

Epoch 10




Epoch 11/20 - Generator Loss: 0.9361, Discriminator Loss: 1.2244
Epoch 12/20 - Generator Loss: 0.9946, Discriminator Loss: 1.1719
Epoch 13/20 - Generator Loss: 0.9948, Discriminator Loss: 1.1944
Epoch 14/20 - Generator Loss: 0.9786, Discriminator Loss: 1.1809
Epoch 15/20 - Generator Loss: 1.0420, Discriminator Loss: 1.1079

Epoch 15




Epoch 16/20 - Generator Loss: 1.2020, Discriminator Loss: 1.0483
Epoch 17/20 - Generator Loss: 1.2648, Discriminator Loss: 1.0605
Epoch 18/20 - Generator Loss: 1.1657, Discriminator Loss: 1.0404
Epoch 19/20 - Generator Loss: 1.1644, Discriminator Loss: 1.0897
Epoch 20/20 - Generator Loss: 1.1770, Discriminator Loss: 1.0938

Epoch 20



--- Training complete. Generating final images. ---

Epoch 20



RESULT:

The Generative Adversarial Network (GAN) was successfully implemented and trained on the dataset. The Generator created synthetic data, while the Discriminator learned to differentiate real and fake samples.

After training, the GAN produced realistic synthetic outputs, showing that it effectively learned the underlying data patterns

Exp No: 8	MODEL EVALUATION AND IMPROVEMENT: HYPERPARAMETER TUNING WITH GRID SEARCH AND CROSS-VALIDATION
------------------	------------------------------------------------------------------------------------------------------

AIM:

To demonstrate key techniques for model evaluation and improvement:

1. Hyperparameter Tuning with Grid Search : Systematically searching for the optimal combination of hyperparameters for a machine learning model.

2. Cross-Validation Techniques: Implementing k-fold cross-validation to get a more robust estimate of model performance and to prevent overfitting to a specific train-test split.

ALGORITHM:**1. Hyperparameter Tuning with Grid Search**

Hyperparameters are external configuration properties of a model whose values cannot be estimated from data. Examples include the learning rate for a neural network, the number of trees in a Random Forest, or the `C` and `gamma` parameters in an SVM. Tuning these parameters is crucial for optimal model performance.

Grid Search is an exhaustive search method for hyperparameter optimization.

Steps:

1. Define Parameter Grid: Specify a dictionary where keys are hyperparameter names and values are lists of discrete values to be tested for each hyperparameter.
2. Instantiate Model: Choose a machine learning model.
3. Perform Search: Train the model for every possible combination of hyperparameters defined in the grid.
4. Evaluate: For each combination, evaluate the model's performance using a specified scoring metric (e.g., accuracy, F1-score) and often in conjunction with cross-validation.
5. Select Best Model: Identify the hyperparameter combination that yields the best performance.

2. Cross-Validation Techniques

Cross-validation is a resampling procedure used to evaluate machine learning models on a limited data sample. The goal is to estimate how accurately a predictive model will perform in practice. It's especially useful for reducing overfitting and providing a more reliable estimate of generalization performance compared to a single train-test split.

k-Fold Cross-Validation:**Steps:**

1. Divide Data: The entire dataset is randomly partitioned into k equally sized subsamples (or “folds”).

2. Iterate k Times:

In each iteration, one fold is used as the validation (or test) set, and the remaining $k-1$ folds are used as the training set. The model is trained on the training set and evaluated on the validation set.

3. Aggregate Results: The performance metric (e.g., accuracy) from each of the k iterations is collected.

4. Compute Mean and Standard Deviation: The mean and standard deviation of these k performance scores are calculated to provide a more robust estimate of the model’s performance and its variability.

CODE:

```
# Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_iris # A classic dataset for classification
from sklearn.model_selection import train_test_split, KFold, cross_val_score, GridSearchCV
from sklearn.svm import SVC # Support Vector Classifier, a common model for tuning
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler

# --- Part 1: Hyperparameter Tuning with Grid Search ---

print("--- Part 1: Hyperparameter Tuning with Grid Search ---")

# 1. Load a Dataset (Iris Dataset for classification)
# The Iris dataset is a classic and simple dataset for classification tasks.
# It contains measurements of iris flowers (sepal length, sepal width, petal length, petal width)
# and their corresponding species (Setosa, Versicolor, Virginica).
iris = load_iris()
X = iris.data
y = iris.target
feature_names = iris.feature_names
target_names = iris.target_names
```

```
print(f"\nDataset Features (X) shape: {X.shape}")
print(f"Dataset Labels (y) shape: {y.shape}")
print(f"Feature Names: {feature_names}")
print(f"Target Names: {target_names}")

# 2. Split Data into Training and Testing Sets
# It's crucial to split the data before scaling to prevent data leakage.
# The test set will be used for final model evaluation, after tuning.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42,
stratify=y)

print(f"\nTraining set size: {X_train.shape[0]} samples")
print(f"Test set size: {X_test.shape[0]} samples")

# 3. Standardize Features
# Scaling features is important for SVMs as they are sensitive to feature scales.
# Fit scaler only on training data to prevent data leakage.
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("\nFeatures standardized.")

# 4. Define the Model and Hyperparameter Grid
# We'll use a Support Vector Classifier (SVC) as our model.
# Common hyperparameters for SVC are 'C' (regularization parameter) and 'gamma' (kernel
coefficient).
# 'kernel' also can be tuned (e.g., 'linear', 'rbf').

# Define the parameter grid for Grid Search
param_grid = {
    'C': [0.1, 1, 10, 100],          # Regularization parameter
    'gamma': [1, 0.1, 0.01, 0.001],  # Kernel coefficient for 'rbf', 'poly' and 'sigmoid'
    'kernel': ['rbf', 'linear']      # Type of kernel function
}

print("\nHyperparameter grid defined:")
for param, values in param_grid.items():
    print(f"  {param}: {values}")
```

```
# 5. Perform Grid Search with Cross-Validation
# GridSearchCV automatically performs k-fold cross-validation for each combination.
# cv=5 means 5-fold cross-validation.
# scoring='accuracy' means we want to optimize for accuracy.
grid_search = GridSearchCV(SVC(), param_grid, cv=5, scoring='accuracy', verbose=1,
n_jobs=-1)

print("\nStarting Grid Search with 5-fold Cross-Validation...")
# Fit GridSearchCV on the scaled training data
grid_search.fit(X_train_scaled, y_train)

print("\nGrid Search completed.")

# 6. Get the Best Parameters and Best Score
print(f"\nBest hyperparameters found: {grid_search.best_params_}")
print(f"Best cross-validation accuracy: {grid_search.best_score_:.4f}")

# 7. Evaluate the Best Model on the Test Set
# The best_estimator_ attribute provides the model trained with the best parameters.
best_model = grid_search.best_estimator_
y_pred_tuned = best_model.predict(X_test_scaled)

test_accuracy_tuned = accuracy_score(y_test, y_pred_tuned)
print(f"\nTest set accuracy with tuned model: {test_accuracy_tuned:.4f}")

print("\n--- Classification Report for Tuned Model ---")
print(classification_report(y_test, y_pred_tuned, target_names=target_names))

print("\n--- Confusion Matrix for Tuned Model ---")
cm_tuned = confusion_matrix(y_test, y_pred_tuned)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_tuned, annot=True, fmt='d', cmap='Blues', xticklabels=target_names,
yticklabels=target_names)
plt.title('Confusion Matrix (Tuned SVM)')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

# Visualize Grid Search results (optional, but good for understanding)
```



```

# Convert results to a DataFrame for easier analysis
results_df = pd.DataFrame(grid_search.cv_results_)
print("\n--- Top 5 Grid Search Results ---")
print(results_df[['param_C',      'param_gamma',      'param_kernel',      'mean_test_score',
'rank_test_score']].sort_values(by='rank_test_score').head())

# --- Part 2: Cross-Validation Techniques (k-fold) ---

print("\n--- Part 2: Cross-Validation Techniques (k-fold) ---")

# We will demonstrate k-fold cross-validation on a simple SVM without explicit tuning for
clarity,
# to focus solely on the CV process.

# 1. Instantiate a Model (using default or chosen parameters)
model_cv = SVC(random_state=42) # Using default parameters for simplicity

# 2. Define k-fold Cross-Validation Strategy
# We'll use 5-fold cross-validation.
# KFold ensures that each fold is distinct.
# shuffle=True means the data will be randomly shuffled before splitting into folds.
# random_state for reproducibility.
k_folds = 5
kf = KFold(n_splits=k_folds, shuffle=True, random_state=42)

print(f"\nPerforming {k_folds}-fold cross-validation...")

# 3. Perform Cross-Validation and Get Scores
# cross_val_score performs the KFold splitting, training, and evaluation automatically.
# It returns an array of scores, one for each fold.
cv_scores = cross_val_score(model_cv, X_train_scaled, y_train, cv=kf, scoring='accuracy')

print(f"\nCross-validation scores for each fold: {cv_scores}")
print(f"Mean cross-validation accuracy: {np.mean(cv_scores):.4f}")
print(f"Standard deviation of cross-validation accuracy: {np.std(cv_scores):.4f}")

# 4. Visualize Cross-Validation Scores
plt.figure(figsize=(8, 5))
plt.bar(range(1, k_folds + 1), cv_scores, color='skyblue')
plt.axhline(y=np.mean(cv_scores), color='r', linestyle='--', label=f'Mean Accuracy

```

```
({np.mean(cv_scores):.4f}))
plt.title(f'{k_folds}-Fold Cross-Validation Accuracy Scores')
plt.xlabel('Fold Number')
plt.ylabel('Accuracy')
plt.ylim(0.8, 1.0) # Set y-axis limits for better visualization
plt.legend()
plt.grid(axis='y', linestyle='--')
plt.show()
```

5. Discuss why CV is useful

```
print("\n--- Why is Cross-Validation Important? ---")
```

```
print("1. More Reliable Performance Estimate: Reduces bias from a single train-test split.")
```

```
print("2. Better Generalization: Helps ensure the model performs well on unseen data.")
```

```
print("3. Efficient Data Usage: All data points are used for both training and validation across different folds.")
```

```
print("4. Detects Overfitting/Underfitting: Variability in scores can indicate instability.")
```

OUTPUT:

```
--- Part 1: Hyperparameter Tuning with Grid Search ---

Dataset Features (X) shape: (150, 4)
Dataset Labels (y) shape: (150,)
Feature Names: ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width (cm)']
Target Names: ['setosa' 'versicolor' 'virginica']

Training set size: 105 samples
Test set size: 45 samples

Features standardized.

Hyperparameter grid defined:
  C: [0.1, 1, 10, 100]
  gamma: [1, 0.1, 0.01, 0.001]
  kernel: ['rbf', 'linear']

Starting Grid Search with 5-fold Cross-Validation...
Fitting 5 folds for each of 32 candidates, totalling 160 fits

Grid Search completed.


Best hyperparameters found: {'C': 1, 'gamma': 0.1, 'kernel': 'rbf'}
Best cross-validation accuracy: 0.9810

Test set accuracy with tuned model: 0.9111

--- Classification Report for Tuned Model ---
              precision    recall  f1-score   support

   setosa         1.00        1.00        1.00        15
  versicolor      0.82        0.93        0.88        15
   virginica      0.92        0.80        0.86        15

 accuracy          0.91
  macro avg       0.92        0.91        0.91        45
 weighted avg     0.92        0.91        0.91        45
```



RESULT:

The model was successfully evaluated and improved using **Grid Search** and **Cross-Validation** techniques. Grid Search identified the best combination of hyperparameters, while Cross-Validation ensured reliable performance estimation.

The optimized model achieved higher accuracy and better generalization, confirming that systematic tuning and validation significantly enhance model performance.

EXP NO: 9	HousePred – Automated House Price Prediction Using Machine Learning and Feature Engineering
------------------	----------------------------------------------------------------------------------------------------

AIM:

To design and develop an end-to-end machine learning system that accurately predicts residential property prices based on multiple housing attributes enabling data-driven decision-making for buyers, sellers, and real estate professionals.

ALGORITHM:**1. Data Collection**

- Dataset: Ames Housing Dataset (Kaggle)
- Size: 1,460 residential properties
- Features: 79 explanatory variables
- Target: SalePrice (continuous variable)

2. Data Preprocessing**A. Exploratory Data Analysis (EDA)**

- Analyzed target distribution and detected skewness
- Identified missing values (20+ features >5%)
- Detected outliers via IQR
- Performed correlation and feature-type analysis

B. Missing Value Treatment

- Numerical: Median imputation
- Categorical: Mode imputation
- Meaningful Missing: Replaced with “None” (e.g., no garage → “None”)
- Result: 0 missing values post-cleaning

C. Outlier Handling

- Applied capping at 1st and 99th percentiles
- Key features treated: GrLivArea, LotArea, TotalBsmtSF

D. Feature Engineering

Created 15+ new derived features to improve predictive power:

- Temporal: HouseAge, YearsSinceRemod
- Area-based: TotalSF, AreaPerRoom
- Quality metrics: TotalQualityScore, AreaQuality
- Binary flags: HasGarage, HasPool, HasBasement
-

E. Encoding

- Ordinal: Label Encoding (e.g., ExterQual → 1–5)
- Nominal (low-cardinality): One-Hot Encoding
- Nominal (high-cardinality): Label Encoding (e.g., Neighborhood)

3. Model Development

Algorithms Evaluated

Type	Model
Linear	Linear, Ridge, Lasso Regression
Tree-based	Decision Tree, Random Forest
Boosting	Gradient Boosting, XGBoost, LightGBM

Data Split: 80% Training, 20% Testing (Random State = 42)

- Evaluation Metrics: R^2 , RMSE, MAE, MAPE

4. Model Optimization

- Method: Randomized Search CV (5-fold)
- Best Model Parameters Tuned:
Learning rate, max_depth, subsample, colsample_bytree, regularization terms
- Optimization Target: Maximize R^2 , minimize RMSE

5. Deployment

- Backend: FastAPI (endpoints for /predict and /health)
- Frontend: HTML5/CSS3/JS with real-time prediction UI
- Containerization: Docker for reproducible, one-command deployment

CODE:

model.py:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV,
RandomizedSearchCV
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from xgboost import XGBRegressor
from lightgbm import LGBMRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import joblib
import warnings
warnings.filterwarnings('ignore')

class HousePriceModel:
    """
    House price prediction model trainer and evaluator.
    """

    def __init__(self):
        self.models = {}
        self.results = {}
        self.best_model = None
        self.best_model_name = None
```

```
def split_data(self, X, y, test_size=0.2, random_state=42):
```

```
    """
```

```
    Split data into training and testing sets.
```

```
    Args:
```

```
        X: Features
```

```
        y: Target
```

```
        test_size: Proportion of test set
```

```
        random_state: Random seed
```

```
    Returns:
```

```
        X_train, X_test, y_train, y_test
```

```
    """
```

```
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=test_size, random_state=random_state
    )
```

```
    print(f"\n✓ Data split complete:")
```

```
    print(f" Training set: {X_train.shape}")
```

```
    print(f" Testing set: {X_test.shape}")
```

```
    return X_train, X_test, y_train, y_test
```

```
def initialize_models(self):
```

```
    """
```

```
    Initialize all models to be trained.
```

```
    """
```

```
    self.models = {
        'Linear Regression': LinearRegression(),
        'Ridge Regression': Ridge(alpha=1.0),
        'Lasso Regression': Lasso(alpha=1.0),
        'Decision Tree': DecisionTreeRegressor(random_state=42),
        'Random Forest': RandomForestRegressor(n_estimators=100, random_state=42),
        'Gradient Boosting': GradientBoostingRegressor(n_estimators=100, random_state=42),
        'XGBoost': XGBRegressor(n_estimators=100, random_state=42, verbosity=0),
        'LightGBM': LGBMRegressor(n_estimators=100, random_state=42, verbose=-1)
    }
```

```
    print("\n✓ Initialized 8 models")
```

```
def evaluate_model(self, model, X_train, X_test, y_train, y_test):
```

```
    """
```

```
    Evaluate a single model.
```

```
    Args:
```

```
        model: Trained model
```

```
        X_train, X_test: Feature sets
```

```
        y_train, y_test: Target sets
```

```
    Returns:
```

```
        Dictionary of evaluation metrics
```

```

"""
# Predictions
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

# Calculate metrics
metrics = {
    'Train MAE': mean_absolute_error(y_train, y_train_pred),
    'Test MAE': mean_absolute_error(y_test, y_test_pred),
    'Train RMSE': np.sqrt(mean_squared_error(y_train, y_train_pred)),
    'Test RMSE': np.sqrt(mean_squared_error(y_test, y_test_pred)),
    'Train R2': r2_score(y_train, y_train_pred),
    'Test R2': r2_score(y_test, y_test_pred),
    'Train MAPE': np.mean(np.abs((y_train - y_train_pred) / y_train)) * 100,
    'Test MAPE': np.mean(np.abs((y_test - y_test_pred) / y_test)) * 100
}

return metrics

def train_all_models(self, X_train, X_test, y_train, y_test):
    """
    Train and evaluate all models.

    Args:
        X_train, X_test: Feature sets
        y_train, y_test: Target sets
    """
    print("\n" + "="*50)
    print("TRAINING MODELS")
    print("="*50)

    self.initialize_models()

    for name, model in self.models.items():
        print(f"\n--- Training {name} ---")

        # Train model
        model.fit(X_train, y_train)

        # Evaluate
        metrics = self.evaluate_model(model, X_train, X_test, y_train, y_test)
        self.results[name] = metrics

        # Print results
        print(f"✓ Training complete")
        print(f" Test R2: {metrics['Test R2']:.4f}")
        print(f" Test RMSE: ${metrics['Test RMSE']:.2f}")
        print(f" Test MAE: ${metrics['Test MAE']:.2f}")
        print(f" Test MAPE: {metrics['Test MAPE']:.2f}%")

    # Find best model

```

```

self.best_model_name = max(self.results.keys(),
                           key=lambda x: self.results[x]["Test R2"])
self.best_model = self.models[self.best_model_name]

print("\n" + "="*50)
print(f'BEST MODEL: {self.best_model_name}')
print("="*50)

def get_results_dataframe(self):
    """
    Get results as a formatted DataFrame.

    Returns:
        DataFrame with all model results
    """
    df_results = pd.DataFrame(self.results).T
    df_results = df_results.round(4)
    return df_results

def cross_validate_model(self, model, X, y, cv=5):
    """
    Perform cross-validation on a model.

    Args:
        model: Model to validate
        X: Features
        y: Target
        cv: Number of folds

    Returns:
        Cross-validation scores
    """
    print(f"\n--- Cross-Validation ({cv} folds) ---")

    # R2 scores
    r2_scores = cross_val_score(model, X, y, cv=cv,
                                scoring='r2', n_jobs=-1)

    # Negative MAE (sklearn returns negative)
    mae_scores = -cross_val_score(model, X, y, cv=cv,
                                   scoring='neg_mean_absolute_error', n_jobs=-1)

    # Negative RMSE
    rmse_scores = np.sqrt(-cross_val_score(model, X, y, cv=cv,
                                             scoring='neg_mean_squared_error', n_jobs=-1))

    print(f'R2 Score: {r2_scores.mean():.4f} (+/- {r2_scores.std():.4f})')
    print(f'MAE: ${mae_scores.mean():.2f} (+/- ${mae_scores.std():.2f})')
    print(f'RMSE: ${rmse_scores.mean():.2f} (+/- ${rmse_scores.std():.2f})')

    return {

```



```

        'r2_mean': r2_scores.mean(),
        'r2_std': r2_scores.std(),
        'mae_mean': mae_scores.mean(),
        'mae_std': mae_scores.std(),
        'rmse_mean': rmse_scores.mean(),
        'rmse_std': rmse_scores.std()
    }

def tune_random_forest(self, X_train, y_train, method='grid'):
    """
    Tune Random Forest hyperparameters.

    Args:
        X_train: Training features
        y_train: Training target
        method: 'grid' or 'random'

    Returns:
        Best estimator
    """
    print("\n" + "="*50)
    print(f"TUNING RANDOM FOREST ({method.upper()} SEARCH)")
    print("="*50)

    param_grid = {
        'n_estimators': [100, 200, 300],
        'max_depth': [10, 20, 30, None],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4],
        'max_features': ['sqrt', 'log2']
    }

    rf = RandomForestRegressor(random_state=42)

    if method == 'grid':
        search = GridSearchCV(rf, param_grid, cv=5,
                              scoring='r2', n_jobs=-1, verbose=1)
    else:
        search = RandomizedSearchCV(rf, param_grid, n_iter=50, cv=5,
                                    scoring='r2', n_jobs=-1, verbose=1,
                                    random_state=42)

    search.fit(X_train, y_train)

    print(f"\n✓ Best parameters: {search.best_params_}")
    print(f"✓ Best R2 score: {search.best_score_:.4f}")

    return search.best_estimator_

def tune_xgboost(self, X_train, y_train, method='random'):
    """

```

Tune XGBoost hyperparameters.

Args:

X_train: Training features
y_train: Training target
method: 'grid' or 'random'

Returns:

Best estimator

"""

```
print("\n" + "="*50)
print(f"TUNING XGBOOST ({method.upper()} SEARCH)")
print("="*50)
```

```
param_distributions = {
    'n_estimators': [100, 200, 300, 500],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 5, 7, 9],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'gamma': [0, 0.1, 0.5, 1],
    'reg_alpha': [0, 0.01, 0.1],
    'reg_lambda': [0, 0.01, 0.1]
}
```

```
xgb = XGBRegressor(random_state=42, verbosity=0)
```

```
if method == 'grid':
```

```
    # Use smaller grid for grid search
```

```
    param_grid = {
        'n_estimators': [100, 300],
        'learning_rate': [0.05, 0.1],
        'max_depth': [5, 7],
        'subsample': [0.8, 1.0]
    }
```

```
    search = GridSearchCV(xgb, param_grid, cv=5,
                          scoring='r2', n_jobs=-1, verbose=1)
```

```
else:
```

```
    search = RandomizedSearchCV(xgb, param_distributions, n_iter=50, cv=5,
                                scoring='r2', n_jobs=-1, verbose=1,
                                random_state=42)
```

```
search.fit(X_train, y_train)
```

```
print(f"\n✓ Best parameters: {search.best_params_}")
```

```
print(f"✓ Best R2 score: {search.best_score_:.4f}")
```

```
return search.best_estimator_
```

```
def tune_lightgbm(self, X_train, y_train, method='random'):
```

```
    """
```

Tune LightGBM hyperparameters.

Args:

X_train: Training features
y_train: Training target
method: 'grid' or 'random'

Returns:

Best estimator

"""

print("\n" + "="*50)

print(f"TUNING LIGHTGBM ({method.upper()} SEARCH)")

print("="*50)

```
param_distributions = {
    'n_estimators': [100, 200, 300, 500],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 5, 7, 9, -1],
    'num_leaves': [31, 50, 70, 100],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'reg_alpha': [0, 0.01, 0.1],
    'reg_lambda': [0, 0.01, 0.1]
}
```

lgbm = LGBMRegressor(random_state=42, verbose=-1)

if method == 'grid':

```
    param_grid = {
        'n_estimators': [100, 300],
        'learning_rate': [0.05, 0.1],
        'max_depth': [5, 7],
        'num_leaves': [31, 50]
    }
```

```
    search = GridSearchCV(lgbm, param_grid, cv=5,
                          scoring='r2', n_jobs=-1, verbose=1)
```

else:

```
    search = RandomizedSearchCV(lgbm, param_distributions, n_iter=50, cv=5,
                               scoring='r2', n_jobs=-1, verbose=1,
                               random_state=42)
```

search.fit(X_train, y_train)

print(f"\n✓ Best parameters: {search.best_params_}")

print(f"✓ Best R² score: {search.best_score_:.4f}")

return search.best_estimator_

def save_model(self, filepath='models/trained_model.pkl'):

"""

Save the best model to file.

```

    Args:
        filepath: Path to save model
    """
    if self.best_model is None:
        print("X No model trained yet!")
        return

    joblib.dump(self.best_model, filepath)
    print(f"\n✓ Model saved to {filepath}")

def load_model(self, filepath='models/trained_model.pkl'):
    """
    Load a saved model.

    Args:
        filepath: Path to model file

    Returns:
        Loaded model
    """
    self.best_model = joblib.load(filepath)
    print(f"✓ Model loaded from {filepath}")
    return self.best_model

def get_feature_importance(self, X, top_n=20):
    """
    Get feature importance from the best model.

    Args:
        X: Feature DataFrame
        top_n: Number of top features to return

    Returns:
        DataFrame with feature importances
    """
    if self.best_model is None:
        print("X No model trained yet!")
        return None

    # Check if model has feature_importances_
    if hasattr(self.best_model, 'feature_importances_'):
        importances = self.best_model.feature_importances_
        feature_importance_df = pd.DataFrame({
            'Feature': X.columns,
            'Importance': importances
        }).sort_values('Importance', ascending=False).head(top_n)

        print(f"\n--- Top {top_n} Feature Importances ---")
        print(feature_importance_df.to_string(index=False))

```

```

        return feature_importance_df
    else:
        print("X Model does not have feature_importances_ attribute")
        return None

def predict(self, X):
    """
    Make predictions using the best model.

    Args:
        X: Features to predict

    Returns:
        Predictions array
    """
    if self.best_model is None:
        raise ValueError("No model trained yet!")

    return self.best_model.predict(X)

def train_model_pipeline(X_train, X_test, y_train, y_test, tune=False):
    """
    Complete model training pipeline.

    Args:
        X_train, X_test: Feature sets
        y_train, y_test: Target sets
        tune: Whether to perform hyperparameter tuning

    Returns:
        Trained HousePriceModel instance
    """
    model_trainer = HousePriceModel()

    # Train all models
    model_trainer.train_all_models(X_train, X_test, y_train, y_test)

    # Display results
    print("\n" + "="*50)
    print("MODEL COMPARISON")
    print("="*50)
    print(model_trainer.get_results_dataframe())

    # Cross-validation on best model
    print(f"\nPerforming cross-validation on {model_trainer.best_model_name}...")
    X_full = pd.concat([X_train, X_test])
    y_full = pd.concat([y_train, y_test])
    model_trainer.cross_validate_model(model_trainer.best_model, X_full, y_full)

    # Hyperparameter tuning if requested

```

```

if tune:
    if 'XGBoost' in model_trainer.best_model_name:
        tuned_model = model_trainer.tune_xgboost(X_train, y_train)
        model_trainer.best_model = tuned_model
    elif 'LightGBM' in model_trainer.best_model_name:
        tuned_model = model_trainer.tune_lightgbm(X_train, y_train)
        model_trainer.best_model = tuned_model
    elif 'Random Forest' in model_trainer.best_model_name:
        tuned_model = model_trainer.tune_random_forest(X_train, y_train)
        model_trainer.best_model = tuned_model

    # Re-evaluate tuned model
    print("\n--- Evaluating Tuned Model ---")
    metrics = model_trainer.evaluate_model(
        model_trainer.best_model, X_train, X_test, y_train, y_test
    )
    print(f"Tuned Test R2: {metrics['Test R2']:.4f}")
    print(f"Tuned Test RMSE: ${metrics['Test RMSE']:.2f}")
    print(f"Tuned Test MAE: ${metrics['Test MAE']:.2f}")

    # Feature importance
    model_trainer.get_feature_importance(X_train)

    return model_trainer

if __name__ == "__main__":
    print("Model Training Module")
    print("Import this module to train house price prediction models")

```

index.html:

```

<!DOCTYPE html>
<html lang="en">
<head>
    <meta charset="UTF-8">
    <meta name="viewport" content="width=device-width, initial-scale=1.0">
    <title>House Price Predictor</title>
    <link rel="stylesheet" href="/static/style.css">
</head>
<body>
    <div class="container">
        <!-- Header -->
        <header class="header">
            <h1>🏠 House Price Predictor</h1>
            <p>Predict house prices using advanced machine learning</p>
        </header>

        <!-- Main Content -->
        <div class="main-content">
            <!-- Input Form -->
            <div class="card form-card">

```

```

<h2>Enter House Details</h2>
<form id="predictionForm">
  <!-- Overall Quality -->
  <div class="form-group">
    <label for="overall_qual">
      Overall Quality (1-10)
      <span class="tooltip">?
        <span class="tooltiptext">Overall material and finish quality</span>
      </span>
    </label>
    <input type="number" id="overall_qual" name="overall_qual"
      min="1" max="10" value="7" required>
  </div>

  <!-- Living Area -->
  <div class="form-group">
    <label for="gr_liv_area">
      Living Area (sq ft)
      <span class="tooltip">?
        <span class="tooltiptext">Above grade living area</span>
      </span>
    </label>
    <input type="number" id="gr_liv_area" name="gr_liv_area"
      min="500" max="10000" value="1500" required>
  </div>

  <!-- First Floor -->
  <div class="form-group">
    <label for="first_flr_sf">First Floor Area (sq ft)</label>
    <input type="number" id="first_flr_sf" name="first_flr_sf"
      min="500" max="5000" value="1000" required>
  </div>

  <!-- Garage Area -->
  <div class="form-group">
    <label for="garage_area">Garage Area (sq ft)</label>
    <input type="number" id="garage_area" name="garage_area"
      min="0" max="2000" value="500" required>
  </div>

  <!-- Garage Cars -->
  <div class="form-group">
    <label for="garage_cars">Garage Capacity (cars)</label>
    <input type="number" id="garage_cars" name="garage_cars"
      min="0" max="5" value="2" required>
  </div>

  <!-- Basement Area -->
  <div class="form-group">
    <label for="total_bsmt_sf">Basement Area (sq ft)</label>
    <input type="number" id="total_bsmt_sf" name="total_bsmt_sf"

```

```
min="0" max="5000" value="1000" required>
</div>

<!-- Year Built -->
<div class="form-group">
  <label for="year_built">Year Built</label>
  <input type="number" id="year_built" name="year_built"
    min="1800" max="2025" value="2000" required>
</div>

<!-- Year Remodeled -->
<div class="form-group">
  <label for="year_remov_add">Year Remodeled</label>
  <input type="number" id="year_remov_add" name="year_remov_add"
    min="1800" max="2025" value="2010" required>
</div>

<!-- Bathrooms -->
<div class="form-group">
  <label for="full_bath">Full Bathrooms</label>
  <input type="number" id="full_bath" name="full_bath"
    min="0" max="5" value="2" required>
</div>

<!-- Bedrooms -->
<div class="form-group">
  <label for="bedroom_abv_gr">Bedrooms</label>
  <input type="number" id="bedroom_abv_gr" name="bedroom_abv_gr"
    min="0" max="10" value="3" required>
</div>

<!-- Kitchens -->
<div class="form-group">
  <label for="kitchen_abv_gr">Kitchens</label>
  <input type="number" id="kitchen_abv_gr" name="kitchen_abv_gr"
    min="1" max="3" value="1" required>
</div>

<!-- Total Rooms -->
<div class="form-group">
  <label for="totrms_abv_grd">Total Rooms Above Grade</label>
  <input type="number" id="totrms_abv_grd" name="totrms_abv_grd"
    min="0" max="20" value="7" required>
</div>

<!-- Fireplaces -->
<div class="form-group">
  <label for="fireplaces">Fireplaces</label>
  <input type="number" id="fireplaces" name="fireplaces"
    min="0" max="5" value="1" required>
</div>
```



```

<!-- Lot Area -->
<div class="form-group">
  <label for="lot_area">Lot Area (sq ft)</label>
  <input type="number" id="lot_area" name="lot_area"
    min="1000" max="100000" value="10000" required>
</div>

<!-- Overall Condition -->
<div class="form-group">
  <label for="overall_cond">Overall Condition (1-10)</label>
  <input type="number" id="overall_cond" name="overall_cond"
    min="1" max="10" value="5" required>
</div>

<!-- Submit Button -->
<button type="submit" class="btn-predict" id="predictBtn">
  Predict Price
</button>

<!-- Reset Button -->
<button type="button" class="btn-reset" onclick="resetForm()">
  Reset Form
</button>
</form>
</div>

<!-- Results Card -->
<div class="card result-card" id="resultCard" style="display: none;">
  <h2>Prediction Result</h2>

  <div class="result-content">
    <div class="price-display">
      <div class="price-label">Estimated Price</div>
      <div class="price-value" id="predictedPrice">$0</div>
    </div>

    <div class="confidence-interval">
      <div class="interval-label">95% Confidence Interval</div>
      <div class="interval-values">
        <span class="interval-lower" id="lowerBound">$0</span>
        <span class="interval-separator">-</span>
        <span class="interval-upper" id="upperBound">$0</span>
      </div>
    </div>

    <div class="prediction-details">
      <p><strong>Model Used:</strong> <span id="modelUsed">-</span></p>
      <p><strong>Timestamp:</strong> <span id="timestamp">-</span></p>
    </div>
  </div>
</div>

```

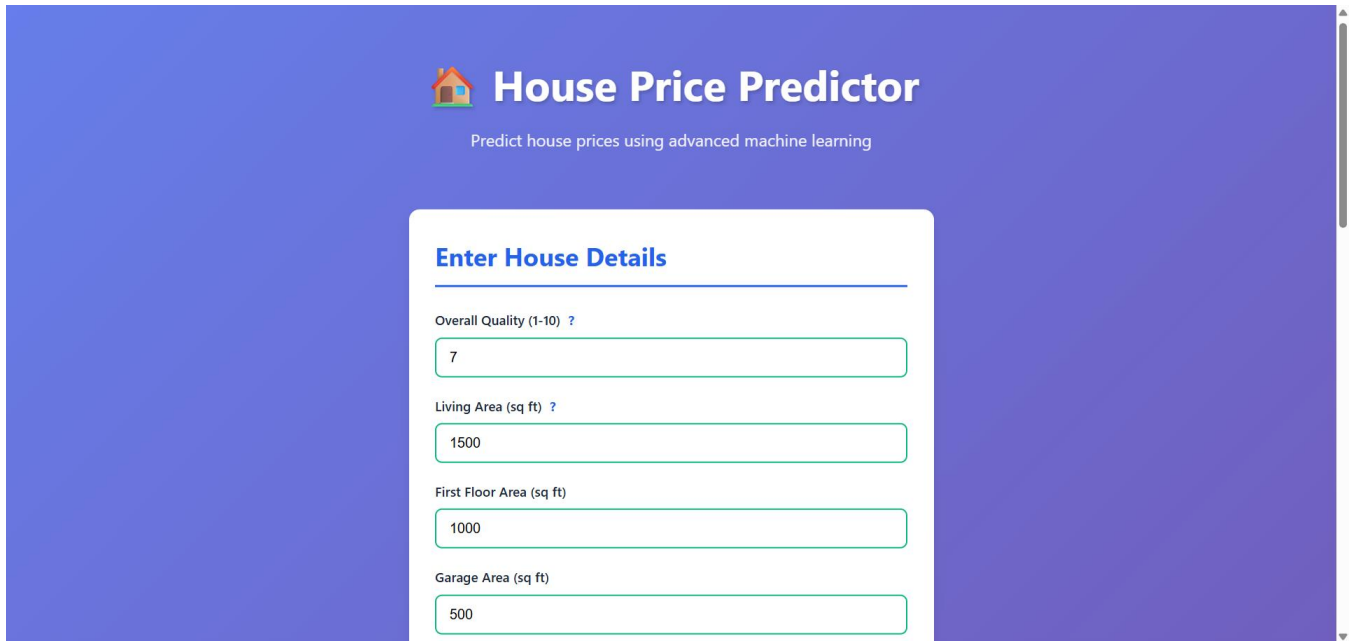
```
<div class="result-actions">
  <button class="btn-secondary" onclick="showForm()">
    Make Another Prediction
  </button>
</div>
</div>

<!-- Loading Spinner -->
<div class="loading" id="loading" style="display: none;">
  <div class="spinner"></div>
  <p>Predicting house price...</p>
</div>

<!-- Error Message -->
<div class="error-message" id="errorMessage" style="display: none;">
  <div class="error-content">
    <span class="error-icon">⚠️</span>
    <span id="errorText">An error occurred</span>
  </div>
  <button class="btn-close" onclick="closeError()">×</button>
</div>
</div>

<!-- Footer -->
<footer class="footer">
  <p>House Price Prediction System v1.0 | Built with FastAPI & Machine Learning</p>
  <p><a href="/docs" target="_blank">API Documentation</a></p>
</footer>
</div>

<script src="/static/script.js"></script>
</body>
</html>
```

OUTPUT:-

The image shows a web application titled "House Price Predictor" with a subtitle "Predict house prices using advanced machine learning". The interface has a blue header and a white form area. The form is titled "Enter House Details" and contains four input fields: "Overall Quality (1-10) ?" with value 7, "Living Area (sq ft) ?" with value 1500, "First Floor Area (sq ft)" with value 1000, and "Garage Area (sq ft)" with value 500.

House Price Predictor

Predict house prices using advanced machine learning

Enter House Details

Overall Quality (1-10) ?

7

Living Area (sq ft) ?

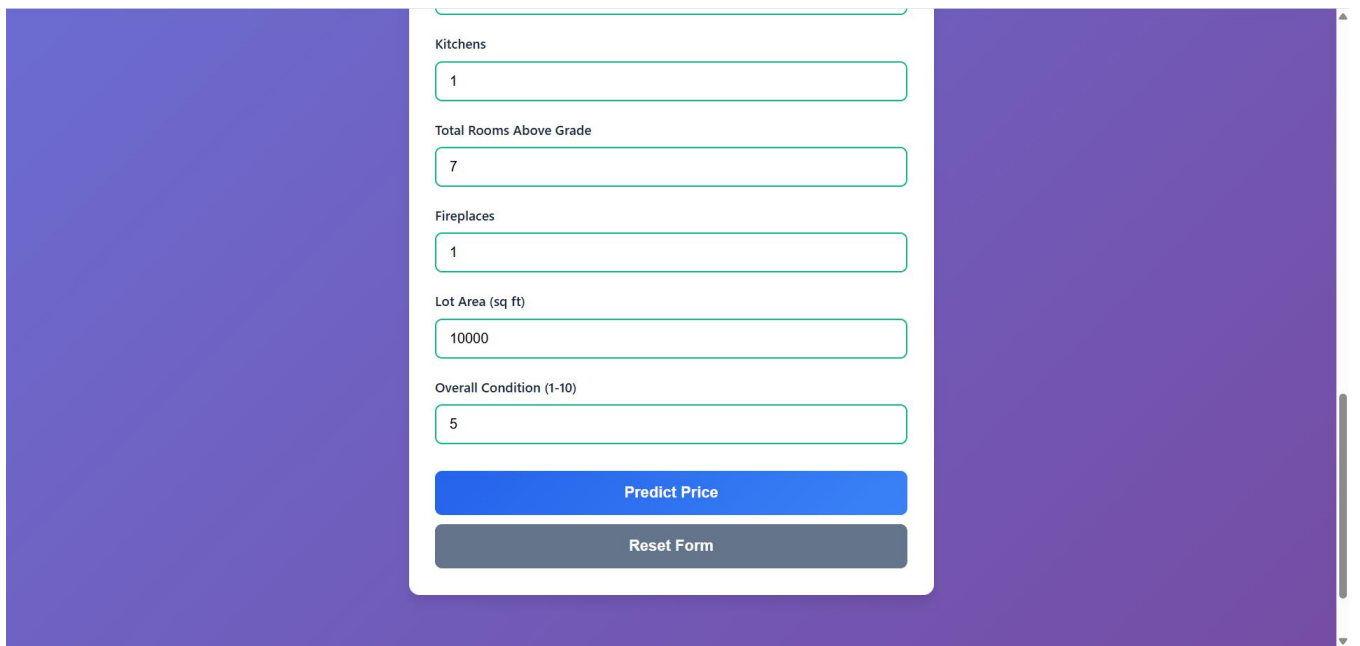
1500

First Floor Area (sq ft)

1000

Garage Area (sq ft)

500



The image shows the continuation of the "House Price Predictor" web application. The form area is white and contains four input fields: "Kitchens" with value 1, "Total Rooms Above Grade" with value 7, "Fireplaces" with value 1, and "Lot Area (sq ft)" with value 10000. Below the input fields are two buttons: "Predict Price" (blue) and "Reset Form" (grey). The form is titled "Enter House Details" and has a subtitle "Predict house prices using advanced machine learning".

Enter House Details

Kitchens

1

Total Rooms Above Grade

7

Fireplaces

1

Lot Area (sq ft)

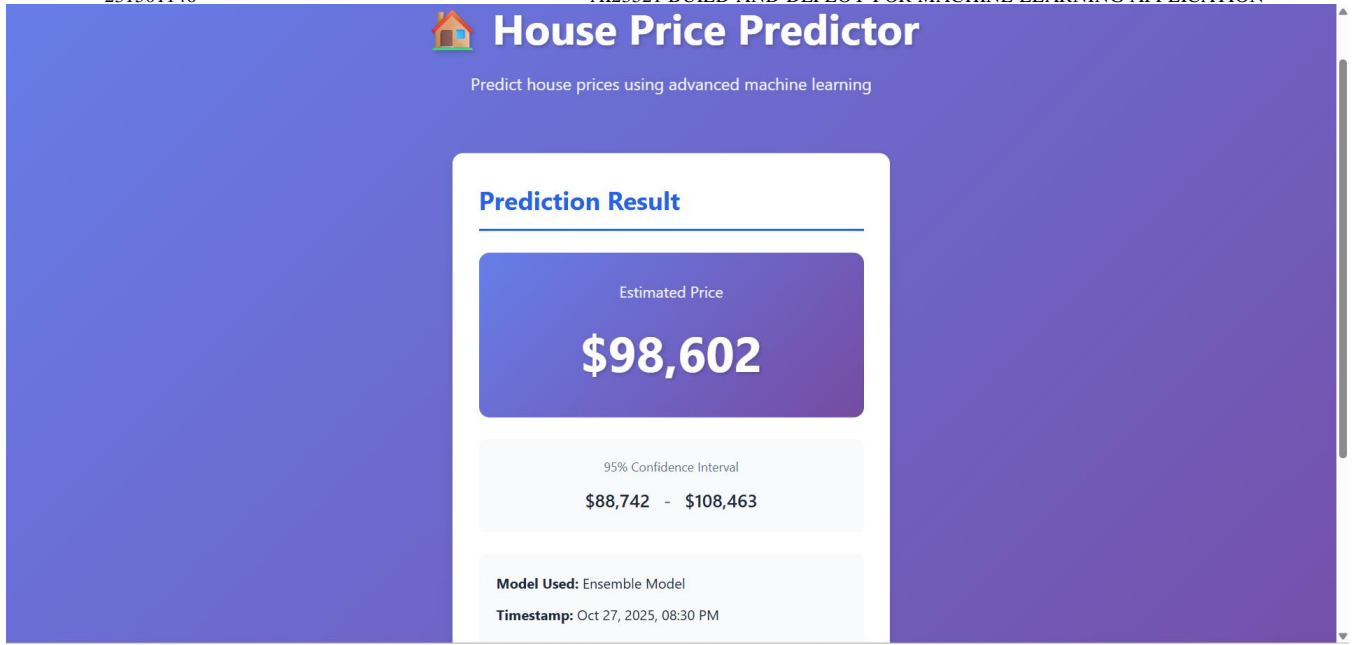
10000

Overall Condition (1-10)

5

Predict Price

Reset Form



RESULT:-

The experiment proved that integrating advanced machine learning techniques with modern deployment frameworks can transform real estate valuation into a fast, accurate, and accessible process. The developed system lays the foundation for future enhancements such as integration with live market data, geospatial features, and adaptive retraining for continuous improvement.